

COMP5338 – Advanced Data Models

Week 1: Big Data, NoSQL and the Polyglot Persistence

Assignment Project Exam Help

Dr. Ying Zhou
School of Information Technologies

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Outline

- The Value and limitations of SQL

- Typical Scale-Out Options

- NoSQL Storage Systems

- Polyglot Persis

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro




The Relational Model of Data

- Entity-Relationship (ER) Data Model describes data as
 - ▶ Entities – distinct objects in the domain
 - ▶ Relationships – between two or more entities
- Entities – described using a set of attributes
- Relationships –
- Used for Conceptual Design
 - ▶ Translated to final database implementation

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



The Rational RDBMS

- Commercial vendors: Oracle, IBM, Microsoft, ...
- Open source systems: MySQL, PostgreSQL, ...
- Common features
 - ▶ Disk-oriented storage
 - ▶ Table stored row
 - ▶ B-trees as index
 - ▶ Dynamic locking
 - ▶ A write-ahead log, or WAL for crash recovery
 - ▶ SQL as the access language
 - ▶ A “row-oriented” query optimizer and executor

Assignment Project Exam Help

<https://eduassistpro.github.io/>

mec

Add WeChat edu_assist_pro



The Value of Relational Databases

- Store persistent data
 - ▶ Storing large amounts of data on disks, while allowing applications to grab the bits they need through queries
- Application Integration
 - ▶ Many applications in an enterprise need to share information, which might happen at the database level
- Concurrency Contr
 - ▶ Database provide nt interaction when many users access the same information at
- Mostly Standard
 - ▶ Relational model is widely used and understood.
 - ▶ SQL is the standard language.
- Reporting
 - ▶ SQL's simple data model and standardization has made it a foundation for many reporting tools

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<http://martinfowler.com/articles/nosql-intro.pdf>



The Scaling Problem of SQL

Relational databases are designed to run on a single machine, so to scale, you need buy a bigger machine or increase capacity of existing server (**scale up**)

But it's cheaper and more effective to **scale out** by buying lots of machines.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<http://martinfowler.com/articles/nosql-intro.pdf>



The Fixed Schema Problem of SQL

- In a relational database
 - ▶ Table structure are predefined
 - ▶ Tables are related with relationships, which are predefined as well
- Schema evolution in RDBMS has large impact on queries and applications
- Example
 - ▶ MediaWiki had versions between April 2003 and November 2007
 - MySQL backend
 - ~ 34 tables, ~242 columns, ~700GB in wikipedia (note: 2008 data)
 - ▶ Schema change has big impact on queries
 - Large number of queries could fail due to schema change.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<http://yellowstone.cs.ucla.edu/schema-evolution/documents/curino-schema-evolution.pdf>



World of Big Data

- Big Data are high-**volume**, high-**velocity**, and/or high-**variety** information assets that require new forms of processing to enable enhanced decision making, insight discovery and process optimization.
[Gartner 2012 report]
- Leaders in database research community identified “big data as a defining challenging of our time” in a 2013 meeting.
- Three major trends
 - “It has become much easier to collect a wide variety of data, due to inexpensive storage, sensors, smart phones, social software, multiplayer games and the Internet of Things, ...”
 - “It has become much cheaper to process large amount data, due to advances in multicore CPUs, solid state storage, inexpensive cloud computing, and open source software”
 - “data management has become democratized. The process of generating, processing, and consuming data is no longer just for database professionals. Decision makers, domain scientists, .. and everyday consumers now routinely do it”

D. Abadi, et al. “The Beckman report on database research”. *Commun. ACM* 59, 2 (January 2016), 92-99



Schema Change is Unavoidable

■ News paper site example

- ▶ Early days, for each news article, we may only record the following information
 - Title, author, publishing date and time, actual content
- ▶ Gradually we may want to record more about an article
 - Keywords, view count, "favoured" it, who emailed it, who twittered it

■ Evolution of an application is in

- ▶ Accept it, incorporate it in the long term of the system
- ▶ Pick a system that allows schema evolution or have a strategy

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Outline

- The Value and limitations of SQL

- **Typical Scale Out Options**

- NoSQL Storage Systems

- Polyglot Persis

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



When Scalability Becomes an Issue

- “**Scalability** is the capability of a system, network, or process to handle a growing amount of work, or its potential to be enlarged in order to accommodate that growth.”

[wikipedia: <https://en.wikipedia.org/wiki/Scalability>]

- In database context, the need for scaling occurs when the **size of the data** **against it** grows to the point of crossing performance
 - ▶ Scale up
 - ▶ Scale out

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Scalability Scenario I

■ Persistent Storage Requirements

- ▶ Medium data size (can fit in one server)
- ▶ Typical query workload consists of **large number of read request** and **relatively low number of write request**

■ Example: wikipedia

- ▶ Only article metadata, user relations, article history, articles stored in core relational database system (MySQL)
- ▶ Article text and images are stored

■ Key challenge

- ▶ Scale to maintain reasonable **read latency**

<http://www.nedworks.org/~mark/presentations/san/Wikimedia%20architecture.pdf>

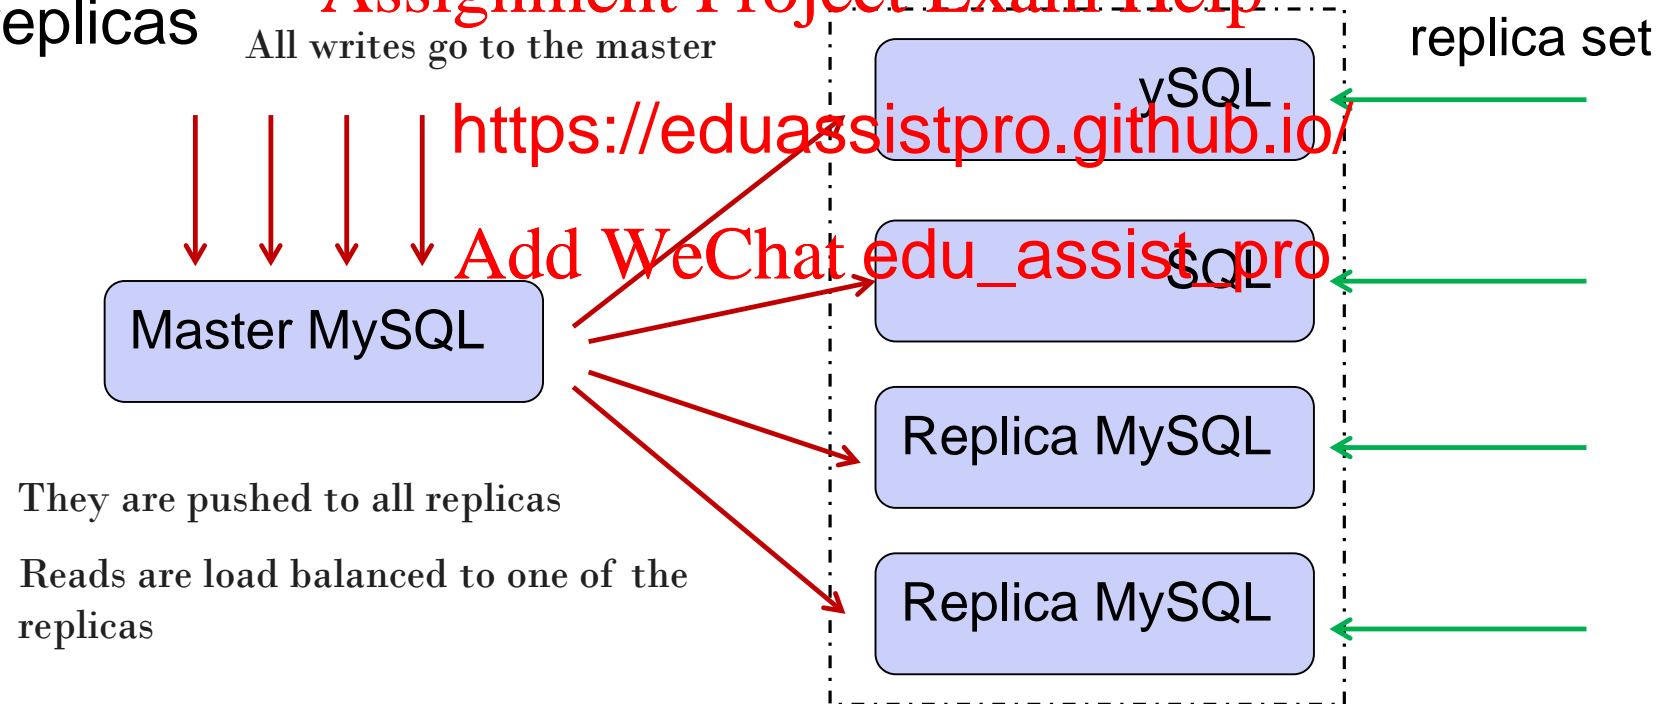


Scalability Scenario I -- Solution

■ Master-Slave Replication

- ▶ Adopted by many companies
- ▶ Also a typical approach to ensure durability

■ Example: Wikipedia has one Master database and many replicas



<http://www.nedworks.org/~mark/presentations/san/Wikimedia%20architecture.pdf>



Scalability Scenario I -- Implications

- When the master dies
 - ▶ One of the replica can be elected as the new master
- Some read may return old data if the latest value has not been pushed from the master
 - ▶ It is possible to let Master handle read request for data requiring strong consistency
- Relatively easy to

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Scalability Scenario II

■ Persistent Storage Requirements

- ▶ Medium or large data size (cannot fit in one server)
- ▶ Typical query workload consists of **large number of read request** and **large number of write request**

■ Example: flickr.com

- ▶ Heavy write traffic, adding favourite, ...
- ▶ Over 400,000 photos (note: 2007 data)
- ▶ More than 4 billion queries per day (note: 2007 data)
- ▶ Uses MySQL as backend storage

■ Key challenge

- ▶ Scale to maintain both **read and write latency**

<http://highscalability.com/flickr-architecture>

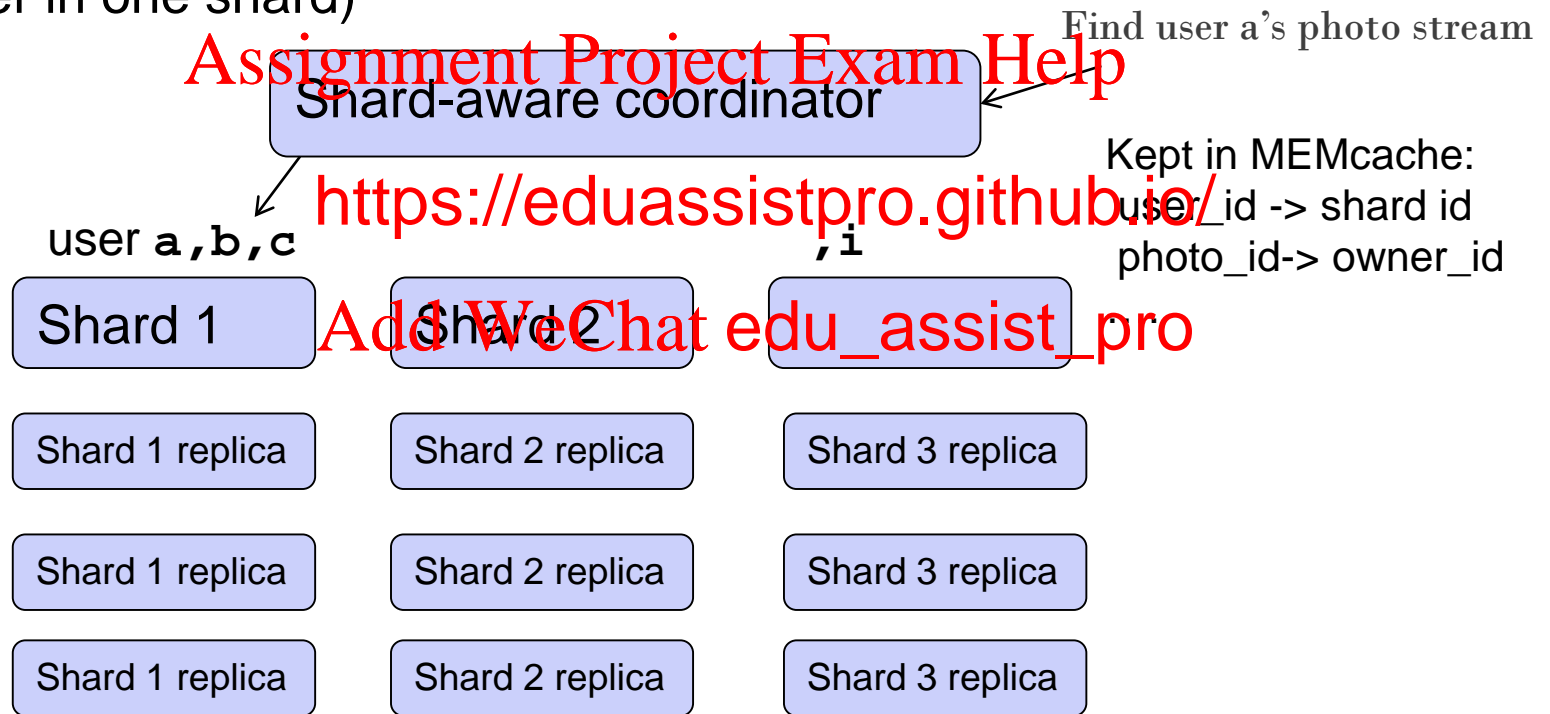
<http://mysqldb.blogspot.com.au/2008/04/mysql-uc-2007-presentation-file.html>



Scalability Scenario II -- Solution

■ Database Sharding

- ▶ The process of slicing a database across multiple machines
- ▶ Most likely horizontally (e.g., store all data related with a particular user in one shard)



Scalability Scenario II – Flickr Example

- User Bob adds User Alice's photo 123 as "favourite"
 - ▶ Pulls the photo (123) owner's account from cache ("Alice"), to get the shard location
 - SHARD-5
 - ▶ Pulls Bob's information from cache, to get Bob's shard location
 - SHARD-13
 - ▶ Start a "distribut

User Bob adds
your photo 123
as favourite

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

User Alice's
photo 123 is
faved by user
Bob

<http://mysqldb.blogspot.com.au/2008/04/mysql-uc-2007-presentation-file.html>



Scalability Scenario II -- Implications

- Data have to be de-normalized
 - ▶ E.g. in the previous example, the “fav” relation is stored in both Bob and Alice’s record.
 - ▶ Join is too expensive when data are sharded
 - ▶ Sometimes join can’t be avoided, e.g. building friends network
- Re-balancing or
 - ▶ What to do whe
- Deciding on a partition factor/pl
 - ▶ May generate hotspots
 - See the twitter example on next slides
- Sharding is largely managed outside RDBMS
 - ▶ Recent version of RDBMS may provide limited support for sharding

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro



Scalability Scenario II – Twitter Example

■ Twitter's problem

- ▶ To store 250 million tweets a day using MySQL

■ Twitter's original Tweet Store:

- ▶ Sharding based on time
- ▶ Range partition (timestamp range)
- ▶ The benefits: sharded sequentially
- ▶ The downsides:
 - Shards with old tweets do not get updated

■ Twitter's new Tweet Store:

- ▶ Sharding based on random partition (id based)
- ▶ A set of in-house systems to manage shards on top of MySQL

<http://highscalability.com/blog/2011/12/19/how-twitter-stores-250-million-tweets-a-day-using-mysql.html>



Outline

- The Value and limitations of SQL

- Handling Scalability

- **NoSQL Storage Systems**

- Polyglot Persis

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



The Coming of NoSQL Storage Systems

- There is no standard definition of NoSQL, the term came up during a workshop on 2009 with presentations from **Voldemort, Cassandra, Dynomite, HBase, Hypertable, CouchDB and MongoDB**
 - ▶ Means “Not Only SQL”
- Typical features
 - ▶ They don't use SQL language thus don't use the SQL language
 - ▶ They don't have fixed schema, all records store any data in any record
 - ▶ Many of them are designed to run on a cluster
 - Manage “sharding”, fault-tolerance, etc. efficiently
 - ▶ Many of them can be integrated with big data processing framework such as MapReduce

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<http://martinfowler.com/articles/nosql-intro.pdf>



NoSQL Ecosystem -- Scalability

■ Distributed NoSQL systems

- ▶ Designed to run on a cluster
- ▶ Support automatically partitioning data across multiple machines
- ▶ Machines can add or leave a *running* cluster
- ▶ Handles failover, fault-tolerance

■ Example Distrib

- ▶ HBase, Cassan

■ Non-distributed NoSQL system

- ▶ Designed to run on a single machine
- ▶ Some has limited support for replication and sharding
- ▶ Schema-less and “object” support

■ Example

- ▶ MongoDB, Neo4j, etc..

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](#)

<http://www.rackspace.com/blog/nosql-ecosystem/>



NoSQL Ecosystem – Data Models

■ Document store

- ▶ Has “table” like concept
- ▶ Each “record” in a “table” is a semi-structured document
- ▶ Examples: MongoDB, CouchDB

■ Column based store

- ▶ Inspired by Google's Bigtable structure
- ▶ Has “table” like concept
- ▶ Storage is organize
- ▶ Examples: Hbase, C

■ Key Value Store

- ▶ Inspired by Amazon's Dynamo storage
- ▶ The overall storage is structured like a big hash table
- ▶ May or may not have a “table” concept
- ▶ Redis, Memcached, Voldemort, S3, Cassandra, DynamnoDB

■ Graph model

- ▶ Storage is organized as nodes and edges
- ▶ Neo4j

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Outline

- The Value and limitations of SQL

- Handling Scalability

- NoSQL Storage Systems

- Polyglot Persis

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Polyglot Persistence

■ SQL is still useful in many ways

- ▶ The relational model fits many real world data
- ▶ ACID transactions are desirable in certain cases
 - Most NoSQL systems have no or very limited transactional support
- ▶ Supporting tools
- ▶ Standard query

■ Polyglot persistence

- ▶ using multiple data storage technologies based upon the way data is being used by individuals.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

large enterprise

Add WeChat edu_assist_pro

<http://martinfowler.com/articles/nosql-intro.pdf>



Data Lake

- Data Lake is a term coined in 2015 referring to a collection of datasets in various formats
 - ▶ Most organizations manage large number of datasets generated or collected from different sources
 - Structured files, databases, spreadsheets, etc
- It also refers to tools that build on top of it.
- There are tools for on house or in the cloud
 - ▶ Google GOODS, Azure Data Lake
- This is still in early stage, with no standard requirement

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Google's GOODS

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Halevy, Alon Y., et al. "Managing Google's data lake: an overview of the Goods system." *IEEE Data Eng. Bull.* 39.3 (2016): 5-14.



Hype Cycle for Data Management, 2017

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<https://www.gartner.com/newsroom/id/3809163>



In Conclusion

- The rise of NoSQL databases marks the end of the era of relational database dominance
- But NoSQL databases will not become the new dominators. Relational will still be popular, and used in the majority of situations. They, however, will no longer be the automatic choice.
- Data lake allows various ways, relational, nonrelational and also as files

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

<http://martinfowler.com/nosql.html>



References

- Paramod J. Sadalage and Martin Fowler, “NoSQL Distilled”, Addison-Wesley Professional; 1 edition (August 18, 2012)
 - ▶ Chapter 1
- Daniel Abadi, Rakesh Agrawal, Anastasia Ailamaki, Magdalena Balazinska, Philip A. Bernstein, Michael J. Carey, Surajit Chaudhuri, Surajit Chaudhuri, Jeffrey Dean, AnHai Doan, Michael J. Franklin, Johannes Gehrke, Laura M. Haas, Alon Y. Halevy, Joseph M. Hellerstein, Yannis E. Ioannidis, H. V. Jagadish, Donald Kossmann, Samuel Madden, Sharad Mehrotra, Tova Milo, Jeffrey F. Naughton, Raghu Ramakrishnan, Volker Markl, Christopher Olston, Beng Chin Ooi, Christopher Ré, Dan Suciu, Michael Stonebraker, and Tom H. Davenport. *Research. Commun. ACM* 59, 2016. doi.org/10.1145/2845915
- Curino, Carlo A., Hyun J. Mo, and Michael J. Carey. “The Evolution of Wikipedia.” *ICEIS*, 2008.
 - ▶ <http://yellowstone.cs.ucla.edu/schema-evolution/document-evolution.pdf>
- High Scalability post: “How Twitter Stores 250 Million Tweets a Day Using MySQL”, 2011
 - ▶ <http://highscalability.com/blog/2011/12/19/how-twitter-stores-250-million-tweets-a-day-using-mysql.html>
- Dathan Pattishall, “Federation at Flickr: Doing Billions of Queries per Day”, 2007
 - ▶ <http://www.scribd.com/doc/2592098/DVPmysqlucFederation-at-Flickr-Doing-Billions-of-Queries-Per-Day>
- Halevy, Alon Y., et al. “Managing Google's data lake: an overview of the Goods system.” *IEEE Data Eng. Bull.* 39.3 (2016): 5-14.
- Gartner Reveals the 2017 Hype Cycle for Data Management
 - ▶ <https://www.gartner.com/newsroom/id/3809163>

