

COMP5338 – Advanced Data Models

Week 3: MongoDB – Aggregation Framework

Assignment Project Exam Help

Dr. Ying Zhou
School of Information Technologies

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



THE UNIVERSITY OF
SYDNEY

Outline

■ Review

■ Aggregation

- ▶ Pipeline stages
- ▶ Operators

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Review

- Document Storage Systems store data as semi-structured document: XML or JSON as two dominant semi-structured formats
 - ▶ Semi-structured data is self-describing
- MongoDB is a popular document storage system that stores data as Binary representation of JSON document (BSON)
- Documents with similar structure representing a particular type of entity are stored in the same collection
- A database is used to represent related entities
- All CRUD operations (find, update, insert, delete) target single collection
- Query criteria, projection and modifier are expressed as JSON document

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Null, empty string and related operators

- Null (or null) is a special data type
 - ▶ Similar to **None**, **Null** or **Nil** in any programming language
 - ▶ It has a singleton value expressed as **null**
 - ▶ Indicating no value is given
- The interpretation of null is different depending on where it appears
- It might represent
 - ▶ The field exists, but has no value
 - ▶ The field does not exist
- This is different to given a field an empty string "" as value

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Null query example

■ Collection revisions document sample

```
{ "_id" : ObjectId("5799843ee2cbe65d76ed919b"),  
  "title" : "Hillary_Clinton",  
  "timestamp" : "2016-07-23T02:02:06Z",  
  "revid" : 7311135,  
  "user" : "BD2412",  
  "parentid" : 7311135,  
  "size" : 251742,  
  "minor" : ""}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- We need a field to indicate if a revision is minor or not. The original schema uses a field with empty string value to indicate a minor revision; a document without this field would be a non-minor revision.

<https://docs.mongodb.com/manual/tutorial/query-for-null-fields/>

Querying for null or field existence

■ Queries

- ▶ `db.revisions.find({minor:{$exists:true}})`
 - Find all documents that a field called `minor` exists
- ▶ `db.revisions.find(where the {minor:""})`
 - Find all documents whose `minor` field has a value of "", empty string
- ▶ `db.revisions.`
 - Find all documents whose `minor` field is null or the value of `minor` field is null
- ▶ `db.revisions.find({minor:{$not:{$exists:true}}})`
 - Find all documents that does not have a field called `minor`

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

It is possible to set the value to null

```
db.revisions.insertOne({title:"nulltest",  
  "timestamp" : "2018-08-14T02:02:06Z",  
  "revid" : NumberLong(7201808141159),  
  "user" : "BD2412",  
  "parentid" : 731113573,  
  "size" : NumberInt(251900),  
  "minor":null})
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
db.revisions.insertOn
```

```
"timestamp" : "2018-08-14T02:02:  
"revid" : NumberLong(201808141157  
"user" : "BD2412",  
"parentid" : NumberLong(731113573),  
"size" : NumberInt(251800))}
```

Add WeChat edu_assist_pro

`db.revisions.find({minor:null})` would return both documents
`db.revisions.find({minor:{$exists:true}})` can differentiate the two



Aggregation

- Simple and relatively standard data analytics can be achieved through **aggregation**

- ▶ Grouping, summing up value, counting, sorting, etc
- ▶ Running on the DB engine instead of application layer

Assignment Project Exam Help

- Several options

- ▶ Aggregation Pip

<https://eduassistpro.github.io/>

- ▶ MapReduce

Add WeChat edu_assist_pro

- Through JavaScript Functions
- Is able to do customized aggregations

Aggregation Pipeline

- Aggregation pipeline consists of multiple stages
 - ▶ Stages are specified using **pipeline operators** such as **\$match**, **\$group**, **\$project**, **\$sort** and so on
 - This is similar to SQL's WHERE, GROUP BY, SORT BY etc
 - Each stage is expressed as an object enclosed by curly bracket
 - ▶ Various **expressions** can be specified in each stage
 - To filter documents based on a condition on an document
 - \$substr, \$
 - ▶ **\$group** stage can specify **accumulator** to perform calculation on documents with the same group key

```
db.collection.aggregate( [
    { pipeline operator: {expression/accumulator,..., expression/accumulator} },
    { pipeline operator: {expression/accumulator,..., expression/accumulator} },
    ...
] )
```

Aggregation Example

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
select cust_id as _id, SUM(amount) as total
from orders
where status = "A"
group by cust_id
```

Typical aggregation stages

- \$match
- \$group
- \$project
- \$sort
- \$skip
- \$limit
- \$count
- \$sample
- \$out
- \$unwind
- \$lookup

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



\$match stage

■ \$match

- ▶ Filter the incoming documents based on given conditions

- ▶ Format:

```
{ $match: { <query> } }
```

- The query document is the same as those in the `find` query

- ▶ Example:

```
db.revisions.aggregate([  
  { $match: { size : { $lt: 250000 } } } ] )
```

Add WeChat edu_assist_pro

Has the same effect as

```
db.revisions.find({ size : { $lt: 250000 } })
```

\$group stage

■ \$group

- ▶ Groups incoming documents by some specified expression and outputs to the next stage a document for each distinct group
 - The `_id` field of the output document has the value of the group key for each group
 - The stage can

```
{ $group: { _id: https://eduassistpro.github.io/  
            <field1>:{accumulator},  
            ...}, Add WeChat edu\_assist\_pro
```

- ▶ To specify the whole collection as a group, give `_id` field `null` value
- ▶ Use *field path* to access fields in the document and set one or many as the value of the `_id` field
 - `"$title"`, or `"$address.street"`
- ▶ There are predefined accumulators: `$sum`, `$avg`, `$first`, `$last`, etc

\$group stage example

- Find the earliest revision time in the whole collection

```
db.revisions.find({}, {timestamp:1, _id:0})
```

```
.sort({timestamp:1})
```

```
.limit(1)
```

```
db.revisions.aggregate([  
  {  
    $group: {  
      _id: null,   
      timestamp: {  
        $min: "$timestamp"   
      }  
    }  
  }  
])
```

- Find the earliest revision time in the collection

```
db.revisions.aggregate([  
  {  
    $group: {  
      _id: "$title",   
      earliest: {  
        $min: "$timestamp"   
      }  
    }  
  }  
])
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

\$group stage example (cont'd)

- Find the number of revisions made on each page by each individual user
 - ▶ This would require grouping based on two fields: title and user
 - ▶ We need to specify these two as the `_id` field of the output document

Assignment Project Exam Help

```
db.revisions.aggregate(  
  {  
    $group: {  
      _id: {  
        title: "$title",  
        user: "$user"  
      },  
      rev_count: {  
        $sum: 1  
      }  
    }  
  }  
])
```

Composite type as `_id`

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

\$group by more than one field

{_id:ObjectId("..."), title: "DT", user:"A", size:123, timestamp:..., ... }
{_id:ObjectId("..."), title: "HC", user:"B", size:113, timestamp:..., ... }
{_id:ObjectId("..."), title: "DT", user:"B", size:125, timestamp:..., ... }
{_id:ObjectId("..."), title: "HC", user:"A", size:113, timestamp:..., ... }
{_id:ObjectId("..."), title: "DT", user:"A", size:125, timestamp:..., ... }

Assignment Project Exam Help

<https://eduassistpro.github.io/>

e:"\$title",user:"\$user"},
rev_count: {\$sum: 1}}}

Add WeChat edu_assist_pro

{_id: {title: "DT", user:"A"}}
{_id: {title: "HC", user:"B"}, rev_count: 1}
{_id: {title: "DT", user:"B"}, rev_count: 1}
{_id: {title: "HC", user:"A"}, rev_count: 1}

\$group examples (cont'd)

- Accumulators do not just return a single value, we can use accumulators to create an array to hold data from incoming documents

- What do the following two commands do:

```
db.revisions.aggregate([
  {$group: {_id:"$title",
    revs: {$push: {"$timestamp"}}}}
])
```

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

```
db.revisions.aggregate([
  {$group: {_id:"$title",rev_users: {$addToSet:"$user"}}}
])
```

\$push accumulator

```
{_id:ObjectId("..."), title: "DT", user:"A", size:123, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "HC", user:"B", size:113, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "DT", user:"B", size:125, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "HC", user:"A", size:113, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "DT", user:"A", size:125, timestamp:..., ... }
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

db.revisions.aggregate([

]])

{user:"A",timestamp:"\$timestamp"}}}

```
{ _id: "DT",
  revs:[
    {user:"A",timestamp:...},
    {user:"B",timestamp:...},
    {user:"A",timestamp:...}
  ]
}
```

```
{ _id:"HC",
  revs:[
    {user:"A", timestamp:...},
    {user:"B", timestamp:...}
  ]
}
```

Add WeChat edu_assist_pro

\$addToSet accumulator

```
{_id:ObjectId("..."), title: "DT", user:"A", size:123, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "HC", user:"B", size:113, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "DT", user:"B", size:125, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "HC", user:"A", size:113, timestamp:..., ... }
```

```
{_id:ObjectId("..."), title: "DT", user:"A", size:125, timestamp:..., ... }
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
ate([  
  title,  
  { $addToSet: "$user" } ] )
```

```
{ _id: "DT",  
  rev_users: [ "A", "B" ]  
}
```

```
{ _id: "HC",  
  rev_users: [ "A", "B" ]  
}
```

\$project stage

■ \$project

- ▶ Reshape the document by including/excluding field, adding new fields, resetting the value of existing field

- ▶ More powerful than the *project* argument in **find** query

- ▶ Format **Assignment Project Exam Help**

{*\$project*: {<sp

- ▶ The specification **<https://eduassistpro.github.io/>** is followed by a single value indicating the inclusion or exclusion
- ▶ Or it can be a field name (existing field) or an expression to compute the value of the field **Add WeChat edu_assist_pro**

<field>: <expression>

- ▶ In the expression, existing field from incoming document can be accessed using field path: “*\$fieldname*”

\$project examples

- Find the age of each title in the collection, where the age is defined as the duration between the last and the first revision of that title, assuming the timestamp is of ISODate type

```
db.revisions.aggregate([
  { $group: { _id: "$title",
    first: { $min: "$timestamp",
    last: { $max: "$timestamp" } },
    age: { $subtract: [ "$last", "$first" ] } } }
])
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

\$group then \$project

```
{_id:ObjectId(""), title: "DT", timestamp:"2016-07-01 00:03:46.000Z", ... }  
{_id:ObjectId(""), title: "HC", timestamp:"2016-07-01 00:55:44.000Z", ... }  
{_id:ObjectId(""), title: "DT", timestamp:"2016-07-15 12:22:35.000Z", ... }  
{_id:ObjectId(""), title: "HC", timestamp:"2016-07-28 00:03:58.000Z", ... }  
{_id:ObjectId(""), title: "DT", timestamp:"2016-07-28 00:20:19.000Z", ... }
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
{ $group: { _id: "$title",  
            { $min: "$timestamp",  
              { $max: "$timestamp" } } },
```

```
{_id:"DT", first:"2016-07-01 00:03:46.000", last:"2016-07-28 00:20:19.000Z"}  
{_id:"HC", first:"2016-07-01 00:55:44.000", last:"2016-07-28 00:03:58.000Z"}
```

Add WeChat edu_assist_pro

```
{ $project: { _id: 0,  
              title: "$_id",  
              age: $subtract: [ "$last", "$first" ] } }
```

```
{title: "DT", age:2333793000}
```

```
{title: "HC", age:2329694000}
```

We can combine multiple operators

```
db.revisions.aggregate([
  {$group: {_id:"$title",
            first: {$min:"$timestamp"},
            last: {$max:"$timestamp"}},
  {$project: {_id:
            tit
            age: {$subtract: [
              [{$subtract: ["$last", "$first"]},
              86400000]}]}
            age_unit: {$literal: "day"}}}]
])
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro



\$sort, \$skip, \$limit and \$count stages

- **\$sort** stage sorts the incoming documents based on specified field(s) in ascending or descending order
 - ▶ The function and format is similar to the sort modifier in **find** query
 - ▶ { **\$sort**: { <field1>: <sort order>, <field2>: <sort order> ... } }
- **\$skip** stage skips over given number of documents
 - ▶ The function and format is similar to the skip modifier in **find** query
 - ▶ { **\$skip**: <positive integer> }
- **\$limit** stage limits the number of documents passed to the next stage
 - ▶ The function and format is similar to the limit modifier in **find** query
 - ▶ { **\$limit**: <positive integer> }
- **\$count** stage counts the number of documents passing to this stage
 - ▶ The function and format is similar to the count modifier in **find** query
 - ▶ { **\$count**: <string> }
 - ▶ String is the name of the field representing the count

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

\$sample and \$out stages

- The **\$sample** stage randomly selects given number of documents from the previous stage
 - ▶ `{ $sample: { size: <positive integer> } }`
 - ▶ Different sampling approaches depending on the location of the stage and the size of the sample and the collection
 - ▶ May fail due to
- The **\$out** stage <https://eduassistpro.github.io/> a given collection
 - ▶ should be the last one in the pipeline
 - ▶ `{ $out: "<output-collection>" }`

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](#)

\$lookup stage

- New aggregation stages are added with major releases.
- **\$lookup** stage is added since 3.2 to perform left outer join between two collections
 - ▶ The collection already in the pipeline (maybe after a few stages)
 - ▶ Another collection (could be the same one)
- For each incoming document the \$lookup stage adds a new **array field** with documents from the other collection.

```
{ $lookup:
  { from: <collection to join>,
    localField: <field from the input documents>,
    foreignField: <field from the documents of the "from"
collection>,
    as: <output array field>
  }
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

\$lookup stage (cont'd)

- The output of **\$lookup** stage has the same number of documents as the previous stage
- Each document is augmented with an **array field** storing matching document(s) from the other collection
- The array could be empty or contain one or more documents depending on the match condition
- Missing local or foreign field is **having null value**

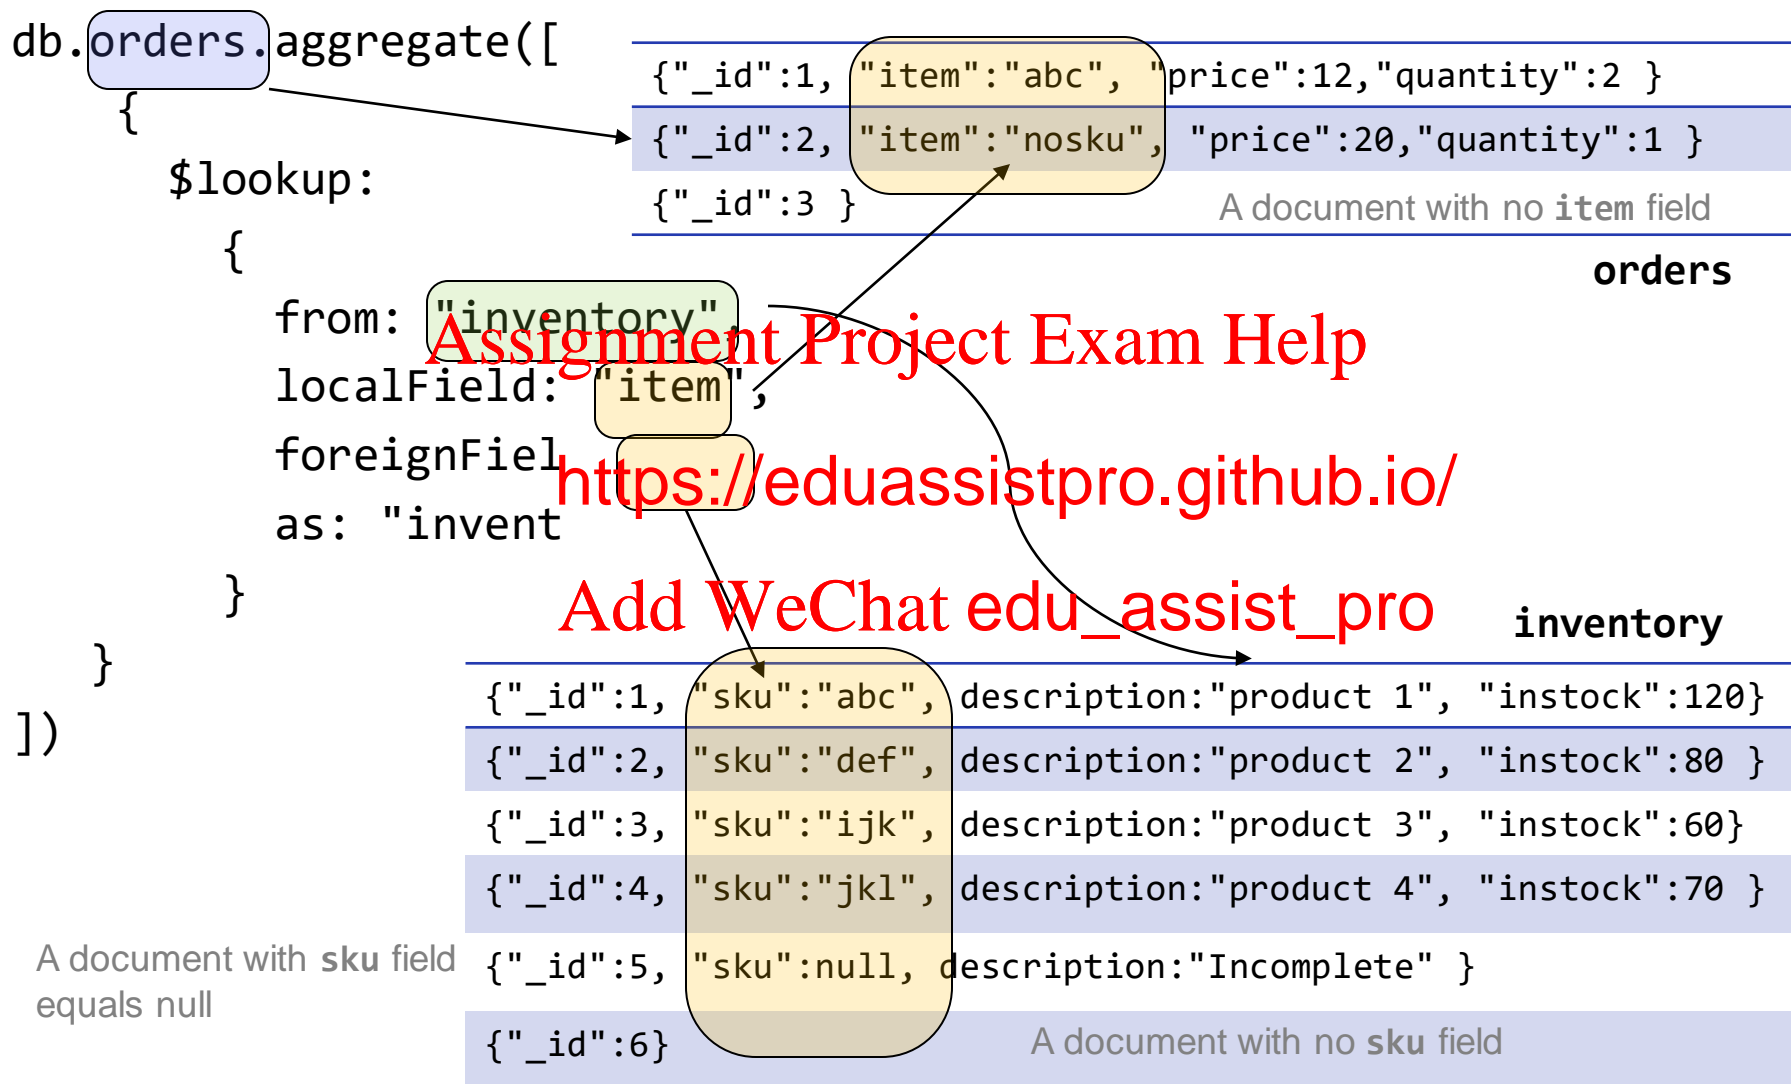
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



\$lookup stage example



https://docs.mongodb.com/manual/reference/operator/aggregation/lookup/#pipe._S_lookup

\$lookup stage example (cont'd)

```
{"_id":1, "item":"abc", "price":12,"quantity":2 }
```

```
{"_id":2, "item":"nosku", "price":20,"quantity":1 }
```

```
{"_id":3 }
```

```
{"_id":1, "sku":"abc", description:"product 1", "instock":120}
```

```
{"_id":2, "sku":"def", description:"product 2", "instock":80 }
```

```
{"_id":3, "sku":"ijk", description:"product 3", "instock":60 }
```

```
{"_id":4, "sku":"jkl", "instock":70 }
```

```
{"_id":5, "sku":null,
```

```
{"_id":6}
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Non-existent null and non exists field

```
{"_id":1, "item":"abc", "price":12,"quantity":2,  
  "inventory_docs": [  
    { "_id":1, "sku":"abc", description:"product 1", "instock":120 } ] }
```

```
{"_id":2, "item":"nosku", "price":20,"quantity":1,  
  "inventory_docs" : [ ] }
```

An empty array for no matching from other collection

```
{"_id":3, "inventory_docs" : [  
  { "_id" : 5, "sku" : null, "description" : "Incomplete" },  
  { "_id" : 6 } ] }
```

Dealing with data of array type

- To aggregate (e.g. grouping) values in an array field, it is possible to flatten the array to access individual value
- **\$unwind** stage flattens an array field from the input documents to output a document for *each* element. Each output document is the input document with the value of the array field replaced by the element.
 - ▶ `{ $unwind: <field path> }`
- Behaviour
 - ▶ Input document:

```
{ "_id" : 1, "item" : "ABC1", "sizes" : [ "S", "M", "L" ] }
```
 - ▶ After **\$unwind**: **"\$sizes"**
 - ▶ Becomes 3 output documents:

```
{ "_id" : 1, "item" : "ABC1", "sizes" : "S" }  
{ "_id" : 1, "item" : "ABC1", "sizes" : "M" }  
{ "_id" : 1, "item" : "ABC1", "sizes" : "L" }
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

\$unwind example

- Find the number of items that are available in each size

```
db.inventory.aggregate( [  
  { $unwind : "$sizes" },  
  { $group: {_id: "$sizes", item_count: {$sum:1}} }  
] )
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



\$unwind then \$group

```
{ "_id" : 1, "item" : "ABC", "sizes": [ "S", "M", "L" ] }  
{ "_id" : 2, "item" : "EFG", "sizes" : [ ] }  
{ "_id" : 3, "item" : "IJK", "sizes": "M" }  
{ "_id" : 4, "item" : "LMN" }  
{ "_id" : 5, "item" : "XYZ", "sizes" : null }
```

Assignment Project Exam Help

{ \$unwind : "\$sizes" },

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

```
{ "_id" : 1  
{ "_id" : 1, "item" : "ABC", "sizes": "M"}  
{ "_id" : 1, "item" : "ABC", "  
{ "_id" : 3, "item" : "IJK", "
```

{ \$group:{_id: "\$sizes",
item_count: {\$sum:1}}

```
{ "_id" : "S", "item_count": 1}  
{ "_id" : "M", "item_count": 2}  
{ "_id" : "L", "item_count": 1}
```


Aggregation Operators

- A few aggregation stages allow us to add new fields or to given existing fields new values based on expression
 - ▶ In `$group` stage we can use various *operators* or *accumulators* to compute values for new fields
 - ▶ In `$project` stage we can use operators to compute values for new or exiting fields
- There are many various data types to carry out common operation
 - ▶ Arithmetic operators: `$mod`, `$subtract`, ...
 - ▶ String operators: `$concat`, `$split`, `$indexOfBytes`, ...
 - ▶ Comparison operators: `$gt`, `$gte`, `$lt`, `$lte`,...
 - ▶ Set operators: `$setEquals`, `$setIntersection`, ...
 - ▶ Boolean operators: `$and`, `$or`, `$not`, ...
 - ▶ Array operators: `$in`, `$size`, ..

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Aggregation vs. Query operators

- There is another set of operators that can be used in **find/update/delete** queries or the **\$match** stage of an aggregation
 - ▶ E.g. **\$gt**, **\$lt**, **\$in**, **\$all**...
- The set is smaller and are different to the operators used in **\$group** or **\$project**
- Some operators have a different syntax and slightly different interpretation in aggregation.
 - ▶ E.g. **\$gt** in query looks like
`{age: {$gt:18}}`
 - ▶ **\$gt** in **\$project** stage looks like:
`{over18: {$gt:["$age", 18]}}`

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Aggregation Behaviour

- It operates on a single collection (before 3.2)
 - ▶ Join can be performed using a particular operator **\$lookup**
- It logically passes the entire collection into the pipeline
- Early filtering can improve the performance
- **\$match** and **\$sort** e index if placed at the beginning of <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Summary

- MongoDB stores data as BSON document
- Retrieving data from MongoDB are usually achieved through
 - ▶ **find** query
 - ▶ **aggregate** pipeline
- **find** query target supports condition on any field
- **aggregate** pipeline is able to process multiple collection(s)
- Both provides rich set of operators
- **update/insert/delete** operation guarantees document level atomicity
- None standard query API, set of operators are growing

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



References

■ BSON types

- ▶ <https://docs.mongodb.com/manual/reference/bson-types/>

■ Aggregation Pipelines

- ▶ <https://docs.mongodb.com/manual/core/aggregation-pipeline/>

■ Aggregation operators

- ▶ <https://docs.mongodb.com/manual/reference/aggregation/>

Assignment Project Exam Help

Add WeChat edu_assist_pro

